

HW 2

Alonso Brito

Problem 1.

10 coins : 3 - 100% Heads , 3 - 100% Tails , 4 - 50/50 H/T

a) ~~$\frac{3}{10} + \frac{3}{10} + \frac{4}{10}(0.5)$~~ $\frac{3}{10} \rightarrow 100\% \text{ of H}$ $\frac{3}{10} \rightarrow 0\%$ $\frac{4}{10} \rightarrow 50\% \text{ of H}$
 $\Rightarrow \frac{3+2}{10} = 0.5 \Rightarrow 50\% \text{ chance of getting a head.}$ $\Rightarrow \frac{4(0.5)=2}{10}$

b) $P(\text{Tail coin} | \text{Tails}) = \frac{P(\text{Tails} | \text{Tail coin}) P(\text{Tail coin})}{P(\text{Tails})}$
 $= \frac{1 (2/3)}{0.5} = \frac{0.3}{0.5} = 0.6 \Rightarrow 60\% \text{ of the coin having Tails on both sides}$
 $P(\text{Fair} | \text{Tails}) = \frac{P(\text{Tails} | \text{Fair}) P(\text{Fair})}{P(\text{Tails})}$
 $= \frac{0.5 (0.4)}{0.5} = 0.4 \Rightarrow 40\% \text{ of the coin being fair}$

c) First coin : 60% of being a Tails coin, 40% of being a fair coin

$\frac{2}{9} \rightarrow 100\% \text{ of T}$	$\frac{3}{9} \rightarrow 100\% \text{ of T}$
$\frac{3}{9} \rightarrow 0\%$	$\frac{3}{9} \rightarrow 0\%$
$\frac{4}{9} \rightarrow 50\% \text{ of T}$	$\frac{3}{9} \rightarrow 50\% \text{ of T}$
$\frac{2}{9} + \frac{2}{9} = \frac{4}{9} = 0.44$	$\frac{3}{9} + \frac{3}{9} = \frac{6}{9} = 50\%$

$(0.44)(0.6) + (0.5)(0.4) = 0.464 \Rightarrow 46.4\% \text{ of getting Tail again}$

Problem 3.

a) $f(x) = \theta \sigma^\theta x^{-\theta-1}$ for $x \geq \sigma = P(x_i | \theta, \sigma)$

$$L(\theta) = \prod_{i=1}^n \theta \sigma^\theta x_i^{-\theta-1}$$

$$l(\theta) = \log L(\theta) = \sum_{i=1}^n \log(\theta \sigma^\theta x_i^{-\theta-1})$$

$$l(\theta) = \sum_{i=1}^n (\log \theta + \theta \log \sigma + (-\theta-1) \log x_i)$$

$$\frac{\partial l(\theta)}{\partial \theta} = \sum_{i=1}^n \left(\frac{1}{\theta} + \log \sigma - \log x_i \right)$$

$$\frac{\partial l(\theta)}{\partial \sigma} = \sum_{i=1}^n \frac{\theta}{\sigma}$$

$$\hat{\theta} = \frac{\partial l(\theta)}{\partial \theta} = 0 \quad \sum_{i=1}^n \left(\frac{1}{\theta} + \log \sigma - \log x_i \right) = 0 = \frac{n}{\theta} + n \log \sigma - \sum_{i=1}^n \log x_i$$

$$\hat{\theta} = \frac{n}{\sum_{i=1}^n (\log x_i - \log \sigma)}$$

$$\hat{\sigma} = \frac{\partial l(\theta)}{\partial \sigma} = 0$$

$$\hat{\sigma} = \sum_{i=1}^n \frac{\theta}{\sigma} = 0$$

$$\hat{\sigma} = \sum_{i=1}^n \theta = n \theta$$

b) $x_i = 13, 2, 16, 5, 11, 16, 18, 5, 8, 15$

$$p(\ln(x_i)) = N(\mu, \sigma^2)$$

$$\hat{\mu} = \sum_{i=1}^n \frac{\ln(x_i)}{n}$$

$$\begin{aligned} \lambda &\Rightarrow \mu \\ \xi &= \sigma \end{aligned}$$

$$\hat{\sigma}^2 = \sum_{i=1}^n \frac{(\ln(x_i) - \mu)^2}{n}$$

from class notes

$$\lambda = 2.709$$

$$\xi = 0.00079$$

c) $f(x; \theta) = \frac{x}{\theta^2} e^{-\frac{x}{\theta}}$ for $x > 0$

$$L(\theta) = \prod_{i=1}^n \frac{x_i}{\theta^2} e^{-\frac{x_i}{\theta}}$$

$$\ell(\theta) = \log L(\theta) = \sum_{i=1}^n \left(\log x_i - \log \theta^2 - \frac{x_i}{\theta} \right)$$

$$\hat{\theta} \Rightarrow \frac{\partial \ell(\theta)}{\partial \theta} = 0$$

$$\frac{\partial \ell(\theta)}{\partial \theta} = \sum_{i=1}^n \left(-\frac{1}{\theta^2} + \frac{x_i}{\theta^3} \right) = \sum_{i=1}^n \frac{1}{\theta^3} (-1 + x_i) = 0$$

~~$$\hat{\theta} = \frac{\sum_{i=1}^n (1 + x_i)}{n}$$~~

$$\hat{\theta} = \sqrt{\sum_{i=1}^n x_i - 1}$$

```
In [1]: import numpy as np
import time

np.random.seed(24787)
a = np.random.randint(8, size=(3,4,4))
```

```
In [2]: print(a)
print('Shape of a: ',a.shape)
```

```
[[[2 6 4 1]
  [0 4 4 3]
  [6 6 1 2]
  [7 0 6 5]]

  [[1 3 3 7]
  [4 7 2 5]
  [0 4 6 7]
  [5 5 7 1]]

  [[7 2 4 5]
  [6 7 7 0]
  [6 2 0 4]
  [2 0 7 6]]]
Shape of a: (3, 4, 4)
```

```
In [3]: fours = np.where(a==4)
#Print indices, first array is depth, second is row, third is column.
print(fours)

(array([0, 0, 0, 1, 1, 2, 2]), array([0, 1, 1, 1, 2, 0, 2]), array([2, 1,
2, 0, 1, 2, 3]))
```

```
In [4]: b = np.tile(a,(2,2))
b.shape
```

```
Out[4]: (3, 8, 8)
```

```
In [5]: c = b.sum(axis=0)

print(c)
print('Shape of c: ',c.shape)
```

```
[[10 11 11 13 10 11 11 13]
 [10 18 13  8 10 18 13  8]
 [12 12  7 13 12 12  7 13]
 [14  5 20 12 14  5 20 12]
 [10 11 11 13 10 11 11 13]
 [10 18 13  8 10 18 13  8]
 [12 12  7 13 12 12  7 13]
 [14  5 20 12 14  5 20 12]]
Shape of c: (8, 8)
```

```
In [6]: np.random.seed(24787)
a = np.random.randint(8, size=(1000,1000))
b = np.random.randint(8, size=(1000,1000))
```



```
In [7]: def matmul(a,b):
        product = np.zeros(a.shape)
        for x in range(a.shape[0]):
            for i in range(a.shape[0]):
                for j in range(a.shape[1]):
                    product[i,x] += a[i,j]*b[j,x]

        return product
```

```
In [8]: start = time.time()
        c = matmul(a,b)
        print("Multiplication took: ", time.time()-start, " seconds.")
```

Multiplication took: 515.9055788516998 seconds.

```
In [9]: start = time.time()
        d = a@b
        print("Multiplication took: ", time.time()-start, " seconds.")
```

Multiplication took: 0.7417197227478027 seconds.

```
In [10]: c==d
```

```
Out[10]: array([[ True,  True,  True, ...,  True,  True,  True],
                [ True,  True,  True, ...,  True,  True,  True],
                [ True,  True,  True, ...,  True,  True,  True],
                ...,
                [ True,  True,  True, ...,  True,  True,  True],
                [ True,  True,  True, ...,  True,  True,  True],
                [ True,  True,  True, ...,  True,  True,  True]])
```

Numpy is generally faster than python because it runs its functions with C and Fortran (static languages), because dynamic languages (like Python) are slow for loops and calling functions.

Question 4: Logistic Regression

```
In [1]: #Import all the required libraries
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
```

Load the data

```
In [2]: # load the data
X1 = pd.read_csv('class0-input.csv')
X2 = pd.read_csv('class1-input.csv')

labels = pd.read_csv('labels.csv')
# Perform important operations on the data
X = pd.concat([X1,X2]).to_numpy()
Y = labels.label.to_numpy()
```

Check the shape


```
In [3]: # Shape of X
print(X.shape)
# Shape of Y
print(Y.shape)
```

```
(10000, 2)
(10000,)
```

Visualize the data

```
In [4]: # Use different colors for each class
# Use plt.scatter
# Dont forget to add axes titles, graph title, legend
plt.scatter(X[Y == 0][:,0], X[Y == 0][:,1], c = 'red')
plt.scatter(X[Y == 1][:,0], X[Y == 1][:,1], c = 'blue')
# plt.scatter(X[5000:,0], X[5000:,1], c = 'blue')
plt.legend(['Class 0', 'Class 1'])
plt.xlabel('x1')
plt.ylabel('x2')
plt.title('Fig. 1: Visualization of X data.')
plt.show()
```

Fig. 1: Visualization of X data.



Define the required functions

```
In [5]: # Pass in the required arguments
# Implement the sigmoid function
def sigmoid(weights, X):
    e = weights[0]
    for i in range(1, len(weights)):
        e += weights[i]*X[:,i-1]
    e = np.exp(e)
    sig = e / (1 + e)
    return sig
```

```
In [6]: # Pass in the required arguments
# The function should return the gradients
def calculate_gradients(pred, X, Y):
    grads = np.array(np.mean(Y - pred))
    for i in range(X.shape[1]):
        grads = np.append(grads, np.mean((Y - pred)*X[:,i]))
    return grads
```

```
In [7]: # Update the weights using gradients calculated using above function and lea
# The function should return the updated weights to be used in the next step
def update_weights(prev_weights, current_grads, learning_rate):

    for i in range(len(prev_weights)):
        prev_weights[i] = prev_weights[i] + learning_rate*current_grads[i]
    return prev_weights
```

```
In [8]: # Use the implemented functions in the main function
# 'main' function should return weights after all the iterations
# Dont forget to divide by the number of datapoints wherever necessary!
# Initialize the intial weights randomly

def main(X, Y, learning_rate = 0.00005, num_steps = 50000):
    np.random.seed(24787)
    weights = np.random.rand(3)
    x1 = X
    for i in range(num_steps):
        sig = sigmoid(weights, X)
        grads = calculate_gradients(sig, X, Y)
        weights = update_weights(weights, grads, learning_rate)
    sig = sigmoid(weights, X)
    grads = calculate_gradients(sig, X, Y)
    weights = update_weights(weights, grads, learning_rate)
    return weights
```

```
In [9]: # Pass in the required arguments (final weights and input)
# The function should return the predictions obtained using sigmoid function
def predict(weights, X):
    Y_hat = sigmoid(weights, X)
    return Y_hat
```

```
In [10]: # Use the final weights to perform prediction using predict function
# Convert the predictions to '0' or '1'
# Calculate the accuracy using predictions and labels

B = main(X,Y)
```

```
In [11]: final_prediction = predict(B,X)


for i in range(len(final_prediction)):
    if final_prediction[i] >= 0.5:
        final_prediction[i] = 1
    else : final_prediction[i] = 0
```

```
In [12]: accuracy = (final_prediction == Y).mean()
print("Accuracy of logistic regression model is: ", accuracy)
```

Accuracy of logistic regression model is: 0.7966

Visualize the misclassification


```
In [13]: # Use different colors for class 0, class 1 and misclassified datapoints
# Use plt.scatter
# Dont forget to add axes titles, graph title, legend
plt.scatter(X[final_prediction == 0][:,0], X[final_prediction == 0][:,1], c
plt.scatter(X[final_prediction == 1][:,0], X[final_prediction == 1][:,1], c
plt.scatter(X[(final_prediction - Y) != 0][:,0], X[(final_prediction - Y) !
plt.legend(['Class 0', 'Class 1', 'Misclassified datapoints'])
plt.xlabel('x1')
plt.ylabel('x2')
plt.title("Fig 2. Visualization of logistic regression model predictions.")
plt.show()
```



Compare the results with sklearn's Logistic Regression

```
In [14]: # import sklearn and necessary libraries
from sklearn.linear_model import LogisticRegression
```


```
In [15]: # Print the accuracy obtained by sklearn and your model
skmodel = LogisticRegression().fit(X,Y)
```

```
In [16]: skmodel.score(X,Y)
```

```
Out[16]: 0.9948
```

```
In [17]: plt.figure(figsize=(10, 6))
plt.scatter(X[Y == 0][:, 0], X[Y == 0][:, 1], color='b', label='0')
plt.scatter(X[Y == 1][:, 0], X[Y == 1][:, 1], color='r', label='1')
plt.legend()
x1_min, x1_max = X[:,0].min(), X[:,0].max(),
x2_min, x2_max = X[:,1].min(), X[:,1].max(),
xx1, xx2 = np.meshgrid(np.linspace(x1_min, x1_max), np.linspace(x2_min, x2_max))
grid = np.c_[xx1.ravel(), xx2.ravel()]
probs = skmodel.predict(grid).reshape(xx1.shape)
plt.contourf(xx1, xx2, probs);
plt.scatter(X[Y == 0][:, 0], X[Y == 0][:, 1], color='b', label='0')
plt.scatter(X[Y == 1][:, 0], X[Y == 1][:, 1], color='r', label='1')
```

Out[17]: <matplotlib.collections.PathCollection at 0x1683e9b10>



With the given hyperparameters, the sklearn model has much better accuracy than my logistic regression model, at 20% higher accuracy score.

Worth noting that, changing the learning rate of my model (removing a 0), its accuracy becomes much closer to the one given by sklearn.